

VISHNUTEJA BIKKINA

Bengaluru, India (Open to relocate: Hyderabad) | +91-8139991104 | bvishnuteja.2004@gmail.com
[linkedin.com/in/vishnutejab](https://www.linkedin.com/in/vishnutejab) | github.com/Vishnuteja156546 | vishnuteja-portfolio.vercel.app

ABOUT ME

AI/ML Developer with hands-on experience building production-grade AI applications using Python, NumPy, Pandas, Matplotlib, RAG pipelines, LLMs, and model fine-tuning. My flagship project is a fully working AutoML desktop application powered by GPT-4 and LangChain – users upload any dataset, chat with an AI assistant, get instant LLM-generated summaries, and receive intelligent chart recommendations. I've trained CNN (92% accuracy) and NLP (86% accuracy) models from scratch, integrated OpenAI APIs into live workflows, and built RAG-based retrieval pipelines for context-grounded generation. Experienced in model evaluation, tuning, and deployment via REST APIs. Available to relocate to Hyderabad immediately.

TECHNICAL SKILLS

LLM & GenAI:

OpenAI API (GPT-4), LangChain, RAG (Retrieval-Augmented Generation), LLM fine-tuning concepts, prompt engineering, AI agent development

Core ML & Deep Learning:

PyTorch, TensorFlow, Keras, Scikit-learn, CNN, RNN, LSTM, model tuning, hyperparameter optimisation, bias-variance tradeoff

NLP:

SpaCy, NLTK, TF-IDF, text classification, summarisation, chatbot development, BLEU, ROUGE, Perplexity

Python & Data:

Python (advanced), NumPy, Pandas, Matplotlib, Seaborn, SQL, ETL pipelines, EDA, feature engineering

Deployment & Tools:

FastAPI, REST APIs, Docker, Streamlit, Tkinter, Git, GitHub, Jupyter Notebook

Analytics & Visualization:

Power BI, Tableau, automated chart recommendation, statistical analysis

WORK EXPERIENCE

AI/ML Developer (Design Engineer)

Dec 2025 – Present

Evolute Fintech, Bengaluru

- Build and deploy AI/ML models using Python, PyTorch, and TensorFlow – covering the full pipeline from data collection (NumPy, Pandas) and EDA (Matplotlib, Seaborn) to model training, tuning, and REST API deployment.
- Integrate OpenAI API and LangChain into business workflows – building LLM-powered summarisation, RAG-based retrieval, and AI copilot applications for fintech data analysis.
- Tune and evaluate models iteratively – applying hyperparameter optimisation, regularisation, dropout, and cross-validation to improve generalisation and production accuracy.
- Automate data pipelines using Python (Pandas, NumPy, SQL) processing 50,000+ records – reducing manual effort by 35% through clean, reproducible ETL workflows.
- Visualise model outputs and insights using Matplotlib, Seaborn, and Power BI – presenting findings clearly to technical and non-technical stakeholders.

AI/ML Intern

Apr 2025 – Sep 2025

Skylkx, Bengaluru

- Built end-to-end NLP and ML pipelines using PyTorch – text preprocessing (SpaCy, NLTK), TF-IDF feature extraction, model training and tuning, evaluated with precision, recall, F1, BLEU.
- Processed 50,000+ records using Pandas and NumPy ETL workflows – delivering clean, model-ready datasets with full documentation.
- Visualised EDA and model results using Matplotlib and Seaborn – creating insight-rich charts for stakeholder reporting.

AI/ML PROJECTS

AutoML AI Application – GPT-4, RAG, LLM, Smart Chart Recommendations *Python, OpenAI API, LangChain, RAG, NumPy, Pandas, Matplotlib, Scikit-learn, Tkinter*

- Built a fully working AutoML desktop application from scratch – users upload any CSV/Excel dataset and the app delivers AI-powered analysis instantly, no code required.
- Integrated GPT-4 via OpenAI API for a conversational AI chatbot – natural language Q&A on uploaded data using RAG pipeline for context-grounded, accurate responses.
- Built LLM-powered summarisation engine – generates plain-English executive summaries covering distributions, correlations, missing values, and trends using Pandas + OpenAI API.
- Implemented intelligent chart recommendation logic using Matplotlib and Seaborn – analyses data types and relationships

to auto-render the most suitable visualisation (bar, scatter, heatmap, histogram, pie).

- Integrated AutoML model selection using Scikit-learn – automatically trains, tunes, and evaluates the best ML model for any uploaded dataset with full metric reporting.
- **GitHub:** github.com/Vishnuteja156546 (add direct repo link)

CNN Image Classifier – Deep Learning, Model Tuning, Deployment

PyTorch, TensorFlow, Keras, NumPy,

Matplotlib, FastAPI, Docker, Streamlit

- Trained CNN from scratch on 5,000+ images – reached 92% accuracy through iterative model tuning: dropout, data augmentation, learning rate scheduling, and batch normalisation.
- Evaluated using precision, recall, F1 – applied bias-variance analysis at each tuning step to balance underfitting and overfitting systematically.
- Deployed via FastAPI REST API + Streamlit frontend, containerised with Docker – full production-ready deployment pipeline.
- **GitHub:** github.com/Vishnuteja156546 (add direct repo link)

NLP Pipeline – LLM-Enhanced Text Classification & Summarisation

PyTorch, SpaCy, NLTK, TF-IDF, OpenAI

API, LangChain, Pandas, FastAPI

- Built complete NLP pipeline – NLTK/SpaCy preprocessing, TF-IDF vectorisation, PyTorch classifier tuned to 86% accuracy, evaluated with BLEU, ROUGE, precision, and recall.
- Extended with OpenAI API for LLM-powered text summarisation and LangChain RAG pipeline for context-aware natural language generation from document corpora.
- Deployed via FastAPI REST endpoint – production-ready, documented, and fast to integrate into any downstream application.

EDUCATION

B.Tech – Computer Science (Data Science)

2021 – 2025

Presidency University, Bengaluru

CERTIFICATIONS

- Python for Data Science – Udemy
- Machine Learning – Simplilearn
- GenAI & LLMs – self-study: OpenAI API, LangChain, RAG pipelines, LLM fine-tuning, prompt engineering